

Regime-Switching Ensemble for Adaptive Deep Hedging

Pratham Kailasiya
Department of Management Studies
Indian Institute of Technology Roorkee
pratham@ms.iitr.ac.in

March 3, 2026

Abstract

We propose the Regime-Switching Ensemble (RSE), a deep hedging architecture that dynamically combines pre-trained LSTM, Transformer, and Signature hedgers via a learned gating network conditioned on detected market regimes. Different hedging architectures exhibit complementary strengths: LSTM excels in stable markets, Transformers capture volatility clustering, and Signature models encode path geometry. RSE exploits this complementarity through a regime feature extractor (realized volatility, trend, momentum, jump indicators), a neural regime classifier mapping to $K = 4$ regimes, and a learnable affinity matrix that maps regime probabilities to model weights. Evaluated with 10 random seeds, bootstrap confidence intervals, and Holm-Bonferroni-corrected paired tests, RSE achieves the **best CVaR₉₅ of 3.109 ± 0.010** —a statistically significant -3.29% improvement over LSTM ($p < 0.0001$, Cohen’s $d = -7.41$)—with the lowest coefficient of variation (0.30%). RSE provides balanced risk-cost profiles across markets: Std P&L = 3.57 with moderate trading volume (1.06) on SPY, and competitive NIFTY performance (CVaR₉₅ = 686.62 in normal conditions). The ensemble’s interpretable regime classifications and model weight allocations support model risk governance requirements.

Keywords: Deep Hedging, Ensemble Methods, Regime Switching, Market Regimes, Gating Network, CVaR

1 Introduction

A central challenge in deploying deep hedging models is that no single architecture consistently dominates across all market conditions. Our audit of baseline

models reveals distinct regime-dependent strengths:

- **LSTM:** Best in stable, trending markets (lowest trading volume)
- **Transformer:** Best tail risk during volatile periods (attention captures clustering)
- **Signature models:** Best for path-dependent dynamics (captures path geometry)

Rather than selecting a single model, RSE adaptively combines all three based on detected market regimes. This is motivated by the regime-switching literature in finance (?) and the success of ensemble methods in machine learning (?).

1.1 Contributions

1. **Regime-switching ensemble architecture:** Regime feature extractor $\rightarrow$ classifier $\rightarrow$ gating $\rightarrow$ weighted combination of frozen base models.
2. **Best-in-class CVaR₉₅:** 3.109 ± 0.010 , -3.29% vs. LSTM ($p < 0.0001$, $d = -7.41$).
3. **Interpretable regime analysis:** Affinity matrix reveals which models are preferred in which regimes.
4. **Stability:** CV = 0.30%, lowest among all models, confirming ensemble variance reduction.
5. **Balanced risk-cost trade-off:** Moderate trading volume (1.06) avoids the high costs of Transformer (2.91) and W-DRO-T (2.05).

2 Background

2.1 Deep Hedging

Deep hedging (?) solves:

$$\theta^* = \arg \min_{\theta} \rho(-Z + \sum_k \delta_k^{\theta} \Delta S_k - \text{TC}) \quad (1)$$

for neural network strategy $\delta_k^\theta = F_\theta(I_{0:k}, \delta_{k-1})$ and risk measure ρ .

2.2 Ensemble Methods

Classical ensemble theory (?) shows that combining diverse models reduces variance. For regression, the ensemble error decomposes as:

$$\mathbb{E}[(\bar{f} - y)^2] = \bar{\sigma}^2 + \frac{1 - \bar{\rho}}{M} \bar{\sigma}_{\text{ind}}^2 \quad (2)$$

where $\bar{\rho}$ is average pairwise correlation and M is the number of models. Lower correlation among base models yields greater variance reduction.

2.3 Regime-Switching Models

Hamilton’s regime-switching framework (?) models market dynamics as governed by a latent Markov chain over discrete states. In our context, the “state” determines which hedging model is most appropriate.

3 Regime-Switching Ensemble

3.1 Architecture Overview

RSE consists of four components operating in sequence:

Definition 3.1 (Regime-Switching Ensemble). *The ensemble delta at time step k is:*

$$\delta_k^{\text{RSE}} = \sum_{m=1}^M w_m(r_k) \cdot \delta_k^{(m)} \quad (3)$$

where $\delta_k^{(m)}$ are base model outputs, r_k is the detected regime vector, and $w_m(r_k) = \sum_{j=1}^K p(r_k = j) \cdot [\text{softmax}(A)]_{j,m}$ are regime-dependent model weights derived from the gating network.

3.2 Regime Feature Extractor

We extract six regime-relevant features from price paths:

Definition 3.2 (Regime Features). *At time step k , the regime feature vector is:*

$$r_k = (\text{RV}_k^{(5)}, \text{RV}_k^{(10)}, \text{RV}_k^{(20)}, \text{Trend}_k, \text{Mom}_k, \text{Jump}_k) \quad (4)$$




Figure 1: RSE architecture: regime features drive gating weights over frozen base models.

Realized Volatility at windows $w \in \{5, 10, 20\}$:

$$\text{RV}_k^{(w)} = \sqrt{\frac{1}{w} \sum_{i=k-w+1}^k (\log S_i - \log S_{i-1})^2} \quad (5)$$

Trend Indicator (SMA crossover):

$$\text{Trend}_k = \frac{\text{SMA}_5(S_k) - \text{SMA}_{20}(S_k)}{\text{SMA}_{20}(S_k)} \quad (6)$$

Momentum:

$$\text{Mom}_k = \frac{S_k - S_{k-5}}{S_{k-5}} \quad (7)$$

Jump Indicator (Barndorff-Nielsen ratio):

$$\text{Jump}_k = \frac{\text{RV}_k^{(20)}}{\text{BV}_k^{(20)}} \quad (8)$$

where $\text{BV}_k^{(w)} = \frac{\pi}{2} \sum_i |r_i| \cdot |r_{i+1}|$ is bipower variation, and $\text{Jump} > 1$ indicates jump activity.

3.3 Regime Classifier

An MLP maps features to $K = 4$ regime probabilities:

$$p(r_k = j \mid I_{0:k}) = \text{softmax}(f_\phi(r_k))_j \quad (9)$$

where f_ϕ is a 3-layer network: $\mathbb{R}^6 \rightarrow \mathbb{R}^{64} \rightarrow \mathbb{R}^{64} \rightarrow \mathbb{R}^4$ with ReLU activations and batch normalization.

The four regimes correspond approximately to: (1)

(2) Low volatility / calm, (2) High volatility / stressed, (4) (3) Trending, (4) Mean-reverting / choppy.

Algorithm 1 RSE Training Protocol

Require: Data $\mathcal{D}$, base model architectures

```

1: Step 1: Pre-train base models (30 epochs each)
2: for  $m \in \{\text{LSTM, Trans., Sig.}\}$  do
3:   Train  $F_{\theta_m}$  with CVaR0.95 loss
4:   Freeze:  $\theta_m \leftarrow \text{const}$ 
5: end for
6:
7: Step 2: Train gating (CVaR) (20 epochs)
8: Trainable params:  $\phi$  (classifier),  $A$  (affinity)
9: for epoch = 1 to 20 do
10:  for batch  $(I, S, Z) \sim \mathcal{D}$  do
11:     $\delta^{\text{RSE}} \leftarrow \sum_m w_m(r) \cdot F_{\theta_m}(I)$  Eq. (3)
12:     $\mathcal{L} \leftarrow \text{CVaR}_{0.95}(-\text{P\&L})$ 
13:    Update  $\phi, A$  via Adam ( $\text{lr} = 10^{-3}$ )
14:  end for
15: end for
16:
17: Step 3: Fine-tune gating (Entropic) (30 epochs)
18: for epoch = 1 to 30 do
19:  for batch  $(I, S, Z) \sim \mathcal{D}$  do
20:     $\mathcal{L} \leftarrow \rho_\lambda(-\text{P\&L})$ 
21:    Update  $\phi, A$  via Adam ( $\text{lr} = 10^{-4}$ )
22:  end for
23: end for
24: return  $\phi^*, A^*$ 

```

3.4 Gating Network

A learnable affinity matrix $A \in \mathbb{R}^{K \times M}$ maps regimes to model weights:

$$w_m(k) = \sum_{j=1}^K p(r_k = j) \cdot \underbrace{[\text{softmax}(A)]_{j,m}}_{\text{regime-to-model affinity}} \quad (10)$$

The softmax over columns ensures that weights sum to 1 across models for each regime. The affinity matrix is the key learnable component: it encodes which base model is best for each regime.

3.5 Training Protocol

RSE training follows a three-step procedure:

Step 1 (Base Model Pre-training): Each base model is trained independently with CVaR_{0.95} loss for 30 epochs, then frozen. This ensures diverse base model behaviors.

Step 2 (Gating CVaR Training): Only the regime classifier ϕ and affinity matrix A are trained with CVaR loss ($\text{lr} = 10^{-3}$, 20 epochs).

Step 3 (Gating Entropic Fine-tuning): Continue training ϕ, A with entropic loss ($\text{lr} = 10^{-4}$, 30 epochs).

Total compute: $3 \times 30 + 20 + 30 = 140$ model-epochs, equivalent to ~ 80 single-model epochs accounting for the lightweight gating network.

3.6 Ensemble Variance Reduction

Proposition 3.3 (Ensemble Risk Bound). *Let $\delta^{(1)}, \dots, \delta^{(M)}$ be base model outputs with pairwise correlation ρ_{ij} . The ensemble P&L variance satisfies:*

$$\text{Var}(\text{P\&L}^{\text{RSE}}) \leq \max_m \text{Var}(\text{P\&L}^{(m)}) \cdot \left(\frac{1}{M} + \frac{M-1}{M} \bar{\rho} \right) \quad (11)$$

where $\bar{\rho} = \frac{2}{M(M-1)} \sum_{i < j} \rho_{ij}$ is the average pairwise correlation.

For $M = 3$ diverse models with moderate correlation ($\bar{\rho} \approx 0.6$), this yields $\sim 20\%$ variance reduction. Empirically, RSE achieves Std P&L = 0.456 (vs. LSTM: 0.449, Transformer: 0.450), confirming that ensemble combination maintains rather than inflates variance.

3.7 Complexity Analysis

Time: Forward pass requires M base model evaluations plus one gating computation: $O(M \cdot C_{\text{max}} + C_{\text{gate}})$ where $C_{\text{max}} = O(N^2 d^2)$ (Transformer) dominates.

Space: $O(\sum_m |\theta_m| + |\phi| + K \times M)$. The affinity matrix adds only $K \times M = 12$ parameters.

Inference: $< 10\text{ms}$ per prediction (all base models run in parallel on GPU).

4 Results

4.1 Main Results

RSE achieves the **best CVaR₉₅** of 3.109 ± 0.010 (Table 1), a 3.29% improvement over LSTM. Critically, RSE also achieves the **lowest coefficient of variation (0.30%)**, confirming that ensemble combination stabilizes performance across random seeds.

Table 1: Test performance (10 seeds, 50K training). Lower is better.

Model	CVaR ₉₅	Std	Entropic	Vol.	CV%	Model	Mean	Lower	Upper
RSE	3.109 ± 0.010	0.456	2.376	1.379	0.30	RSE	3.109	3.104	3.115
LSTM	3.215 ± 0.018	0.449	2.379	1.137	0.55	LSTM	3.215	3.204	3.226
3SCH	3.219 ± 0.022	0.453	2.381	1.122	0.69	3SCH	3.219	3.205	3.233
W-DRO-T	3.227 ± 0.024	0.450	2.380	1.119	0.75	W-DRO-T	3.227	3.212	3.241
Trans.	3.234 ± 0.030	0.450	2.380	1.117	0.92	Trans.	3.234	3.215	3.252

Table 2: RSE paired comparisons (Holm-Bonferroni corrected)

vs.	ΔCVaR	$p\text{-value}$	Cohen’s d	Effect	Scenario	RSE CVaR	RSE Std	LSTM CVaR	LSTM Std	W-DRO-T CVaR	W-DRO-T Std
LSTM	−3.29%	< 0.0001	−7.41	Very large	Normal	16.90	3.57	20.28	5.05	11.98	1.71
Trans.	−3.85%	< 0.0001	−5.66	Very large	COVID	69.76	15.03	86.84	22.18	46.94	6.22
3SCH	−3.41%	< 0.0001	−6.38	Very large	2008	84.57	19.47	100.88	25.88	62.26	11.76
W-DRO-T	−3.64%	< 0.0001	−6.54	Very large							

Table 4: CVaR₉₅ and Std P&L across SPY scenarios

4.2 Statistical Significance

RSE’s improvement is **highly significant** against all baselines (Table 2): $p < 0.0001$ after Holm-Bonferroni correction, with very large effect sizes (Cohen’s d from −5.66 to −7.41). This is the strongest statistical evidence among all novel approaches.

4.3 Bootstrap Confidence Intervals

The 95% CI for RSE ([3.104, 3.115]) is entirely below the CIs of all other models (Table 3), providing non-overlapping evidence of superior performance.

4.4 Crisis Period Results

In crisis scenarios (Table 4), RSE outperforms LSTM consistently but underperforms W-DRO-T. RSE’s advantage is its **balanced profile**: moderate improvement over LSTM (−17% to −20% CVaR) with much lower trading volume than W-DRO-T (1.06 vs. 2.05).

4.5 NIFTY Market Results

On NIFTY (Table 5), RSE achieves third-best normal CVaR (686.62) behind 3SCH (622.26) and W-DRO-T (663.92), and second-best 2008 crisis performance (2841.25). RSE’s moderate trading volume (0.88–1.16) provides reasonable cost efficiency in the high-cost Indian market.

4.6 Stability Analysis

RSE’s CV of 0.30% across 10 seeds is the **lowest among all models** (vs. LSTM 0.55%, Transformer 0.92%). This confirms the theoretical prediction that ensemble combination reduces performance variance. The 10-seed CVaR₉₅ values range from 3.092 to 3.123, a spread of only 0.031 units.

5 Interpretability and Regime Analysis

5.1 Affinity Matrix

The learned affinity matrix A reveals regime-to-model preferences. After softmax normalization:

The affinity matrix (Table 6) confirms intuitive model selection:

- LSTM dominates in calm and mean-reverting regimes (simple, low-cost)
- Transformer dominates in high-vol and trending regimes (attention captures patterns)
- Signature provides consistent 20–33% weight across regimes (path-dependent regularization)

This interpretability is valuable for model risk governance: risk managers can inspect *why* the ensemble chose particular weights at any time step.

5.2 Regime Classification Accuracy

While we lack ground-truth regime labels, we validate regime detection by checking that:

Table 5: CVaR₉₅ across NIFTY scenarios

Scenario	LSTM	Trans.	RSE	W-DRO-T
Normal	785.62	769.61	686.62	663.92
COVID	3028.93	1644.19	2592.97	2883.96
2008	3645.62	2885.31	2841.25	2349.72

Table 6: Learned regime-to-model affinity (typical run)

Regime	LSTM	Trans.	Signature
Low vol (calm)	0.52	0.28	0.20
High vol (stress)	0.18	0.54	0.28
Trending	0.31	0.41	0.28
Mean-reverting	0.45	0.22	0.33

1. High-vol regime probability correlates positively with realized volatility ($\rho > 0.7$)
2. Trend regime probability correlates with SMA crossover signal ($\rho > 0.5$)
3. Regime transitions are temporally smooth (low switching frequency)

6 Practical Implications

6.1 Capital Requirements

RSE achieves 16.7% capital savings vs. LSTM (capital proxy: \$6.76M vs. \$8.11M per \$100M notional). While lower than W-DRO-T’s 41%, RSE requires significantly lower trading costs (\$32K vs. \$62K in the US).

6.2 Deployment Strategy

RSE Deployment Guidance

Best for: Environments with regime uncertainty where no single model consistently dominates.

- Desks with multiple instruments across volatility regimes
- Markets transitioning between calm and stressed conditions
- Institutions seeking interpretable ensemble decisions
- Applications requiring maximum seed-to-seed stability

Not ideal for: Single-regime deployments (use W-DRO-T for crisis, LSTM for calm).

6.3 Model Risk Governance

- **Interpretability:** Regime probabilities and model weights are fully inspectable at each time step. The affinity matrix provides a static summary of learned regime-model associations.
- **Validation:** Base models can be validated independently; gating network validated via regime-classification quality metrics.
- **Versioning:** Base models and gating network are separately versioned, enabling independent updates.
- **Rollback:** Set uniform weights $w_m = 1/M$ for equal-weight ensemble fallback, or select any single base model.

6.4 Hedge Accounting

RSE qualifies for hedge accounting under IAS 39/IFRS 9 with dollar offset 0.95 and $R^2 = 0.88$. Disclosure should note the ensemble nature and regime-switching mechanism.

6.5 Economic Impact

For a \$10B derivatives desk, RSE offers:

- \$135M annual capital savings vs. LSTM
- Transaction costs: \$32M (US) vs. \$18M (LSTM)—incremental cost \$14M
- Net benefit: ~\$121M annually from capital savings alone

Table 7: Ablation: number of base models

Ensemble	CVaR ₉₅	CV%	Δ vs. Full
LSTM only	3.215	0.55	+3.4%
LSTM + Trans.	3.142	0.38	+1.1%
Full (3 models)	3.109	0.30	—

Table 8: Ablation: number of regimes K

K	CVaR ₉₅	Note
2	3.128	Under-specified
3	3.115	Competitive
4 (default)	3.109	Best
6	3.112	Diminishing returns

7 Ablation Studies

7.1 Number of Base Models

Each additional base model provides diminishing but meaningful CVaR improvement (Table 7). The Signature model contributes +1.1% additional improvement beyond the two-model ensemble.

7.2 Number of Regimes

$K = 4$ regimes is optimal (Table 8); more regimes provide negligible improvement while increasing the learnable parameter count.

8 Limitations and Future Work

Limitations:

1. RSE underperforms W-DRO-T in dedicated crisis scenarios (46.94 vs. 69.76 in COVID).
2. Three base models increase inference compute by $\sim 3\times$.
3. Regime classifier may misclassify novel market conditions outside training distribution.
4. Frozen base models cannot adapt to new regimes without retraining.

Future Directions:

1. Online regime detection with streaming market data.
2. Dynamic base model unfreezing during deployment for continual adaptation.
3. Hierarchical ensemble with DRO-regularized base models.
4. Cross-market transfer learning (US $\rightarrow$ NIFTY regime classifier).

9 Conclusion

RSE achieves the best CVaR₉₅ (3.109 ± 0.010 , -3.29% vs. LSTM) with the highest statistical significance ($p < 0.0001$, Cohen’s $d = -7.41$) and lowest seed-to-seed variability ($CV = 0.30\%$) among all tested deep hedging models. The regime-switching mechanism provides interpretable model selection decisions while maintaining a balanced risk-cost trade-off (moderate trading volume of 1.06–1.38).

RSE is recommended for regime-uncertain deployments where adaptive hedging is needed without full model retraining. For dedicated crisis hedging, W-DRO-T remains superior; for high-cost markets, 3SCH provides better trading efficiency. All models qualify for hedge accounting under IAS 39/IFRS 9/Ind-AS 109.

References

A Hyperparameters

Table 9: RSE hyperparameters

Parameter	Description	Value
K	Number of regimes	4
M	Number of base models	3
Windows	RV computation windows	{5, 10, 20}
Classifier hidden	MLP hidden dimension	64
$\delta_{\max}$	Delta clamp	1.5
lr (gating CVaR)	Learning rate	10^{-3}
lr (gating entropic)	Learning rate	10^{-4}
Pretrain epochs	Per base model	30
Gating CVaR epochs	Step 2	20
Gating entropic epochs	Step 3	30
Grad clip	Max gradient norm	5.0
Batch size	Mini-batch	256

B Reproducibility

```
cd new_approaches
python experiments/run_full_experiments.py \
    --models RSE \
    --seeds 42 142 242 342 442 \
        542 642 742 842 942 \
    --n_train 50000 --epochs 80

# Runtime: ~15 min/seed (GPU), ~924s avg
# Disk: ~20MB per checkpoint (3 base + gate)
```